



MONTANA
STATE UNIVERSITY

MongoDB

...

More NoSQL: Document Databases

MongoDB

- Like Redis, MongoDB is a *NoSQL* data store
- Unlike Redis, MongoDB is a *Document Database*
 - Stores JSON-like documents rather than offering various data types



mongoDB®

Documents In MongoDB

- At right is a MongoDB document
- As you can see, it looks a lot like JSON
- Additionally, you can see that field values need not be a primitive type
 - Arrays
 - Nested Objects
 - etc.

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```



← field: value
← field: value
← field: value
← field: value

Documents In MongoDB

- Advantages of this approach
 - Documents (i.e. objects) correspond to native data types in many programming languages
 - Embedded documents and arrays reduce need for expensive joins
 - Dynamic schema supports fluent polymorphism

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

Documents In MongoDB

- *“Documents correspond to native data types in many programming languages”*
 - To an extent
 - Obviously there is a direct mapping to Javascript, but what about other languages?
 - Yeah, but not as good a match
 - Static typing?

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```



← field: value

← field: value

← field: value

← field: value

Documents In MongoDB

- *“Embedded documents and arrays reduce need for expensive joins”*
 - True, to an extent
 - But what about normalization?
 - What if I want to rename the “sports” group to “college sports”?

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

Documents In MongoDB

- “Dynamic schema supports fluent polymorphism”
 - Uhhhh, OK
 - Documents don't need to meet any schema
 - So two documents in the same collection can look totally different...
 - Uhhhh, I guess that works for javascript...

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

Documents In MongoDB

- Mongo 3.2 added support for document validation

- Allows you to enforce a schema via JSONSchema

<https://json-schema.org/>

- Side bar, guess who wrote this:

<http://jschema.org>

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}
```

← field: value
← field: value
← field: value
← field: value

MongoDB Concepts

- Database: a set of named *Collections*
 - As well as views and a few other things
- Collections: a set of Documents
 - Stored in BSON format

<http://bsonspec.org/>

```
        \x31\x00\x00\x00
        \x04BSON\x00
        \x26\x00\x00\x00
        \x02\x30\x00\x08\x00\x00\x00awesome\x00
        \x01\x31\x00\x33\x33\x33\x33\x33\x14\x40
        \x10\x32\x00\xc2\x07\x00\x00
        \x00
        \x00
```

{"BSON": ["awesome", 5.05, 1986]} →

Creating Collections

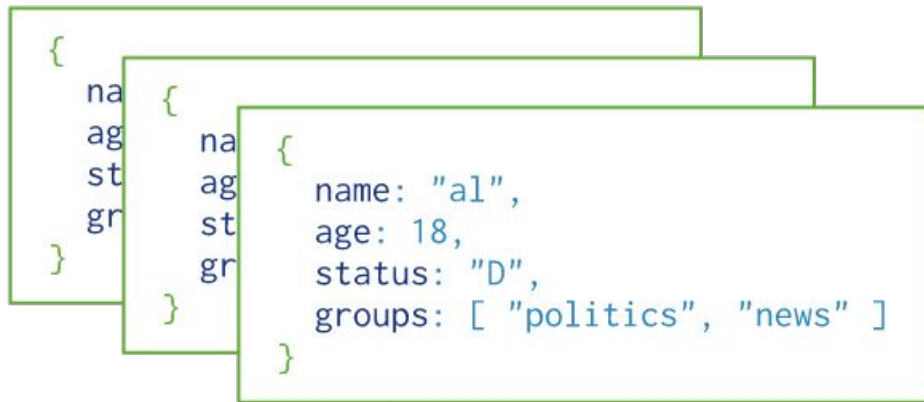
- One really nice feature of Mongo is how easy it is to create a database
- This will create a new database, *myNewDB* and then create a new collection and insert some data into it
 - Very liberating if you are used to DDL

```
use myNewDB
```

```
db.myNewCollection1.insertOne( { x: 1 } )
```

Collections

- Inserting multiple documents into a collection is similar, although not identical, to inserting rows in a relational database



Collection

Collections

- You can explicitly create a collection if you wish to specify configuration options for it
 - Note the validator option
 - You can also cap the size of the collection, etc.

```
db.createCollection( <name>,  
  {  
    capped: <boolean>,  
    autoIndexId: <boolean>,  
    size: <number>,  
    max: <number>,  
    storageEngine: <document>,  
    validator: <document>,  
    validationLevel: <string>,  
    validationAction: <string>,  
    indexOptionDefaults: <document>,  
    viewOn: <string>,           // Added in MongoDB 3.4  
    pipeline: <pipeline>,      // Added in MongoDB 3.4  
    collation: <document>,      // Added in MongoDB 3.4  
    writeConcern: <document>  
  }  
)
```

Views

- As with relational databases, you can create views in MongoDB
- Here `<source>` is a query that will define the view
 - We will discuss the pipeline later
- As with relational database, views are read only

```
db.createView(  
  "<viewName>",  
  "<source>",  
  [<pipeline>],  
  {  
    "collation" : { <collation> }  
  }  
)
```

Documents

- In MongoDB documents you will always have an `_id` field, with is of type `ObjectId`
 - ObjectIds are small, *likely unique*, fast to generate, and ordered. ObjectId values are 12 bytes in length, consisting of:
 - a 4-byte timestamp value, representing the ObjectId's creation, measured in seconds since the Unix epoch
 - a 5-byte random value
 - a 3-byte incrementing counter, initialized to a random value

```
var mydoc = {  
  _id: ObjectId("5099803df3f4948bd2f98391"),  
  name: { first: "Alan", last: "Turing" },  
  birth: new Date('Jun 23, 1912'),  
  death: new Date('Jun 07, 1954'),  
  contribs: [ "Turing machine", "Turing test", "Turingery" ],  
  views : NumberLong(1250000)  
}
```

Documents

- This document also shows
 - nested objects (name)
 - Dates
 - Arrays
 - A NumberLong(64 bit integer)
 - By default the mongodb shell treats numbers as double precision floating point, like javascript
 - `_(ツ)_/`

```
var mydoc = {  
  _id: ObjectId("5099803df3f4948bd2f98391"),  
  name: { first: "Alan", last: "Turing" },  
  birth: new Date('Jun 23, 1912'),  
  death: new Date('Jun 07, 1954'),  
  contribs: [ "Turing machine", "Turing test", "Turingery" ],  
  views : NumberLong(1250000)  
}
```


CRUD in Mongo

- Documents can be created with the *insertOne()* and *insertMany()* operations
 - This example inserts a single value into the example collection
 - Recall that the collection will automatically created if it does not exist
 - Note that you get back the object ID of the inserted document



```
> use demodb
switched to db demodb
> db.example.insertOne({"foo":"bar"})
{
  "acknowledged" : true,
  "insertedId" : ObjectId("5f984793663182f750282296")
}
> db.example.find()
{ "id" : ObjectId("5f984793663182f750282296"), "foo" : "bar" }
```

A screenshot of the MongoDB shell interface. The terminal shows a sequence of commands: `use demodb`, `switched to db demodb`, `db.example.insertOne({"foo":"bar"})`, and `db.example.find()`. The output of the `insertOne` command is a JSON object with `"acknowledged" : true` and `"insertedId" : ObjectId("5f984793663182f750282296")`. The output of the `find` command is a document with `"id" : ObjectId("5f984793663182f750282296")` and `"foo" : "bar"`. A red arrow points to the `insertOne` command line.

CRUD in Mongo

- Querying in Mongo is usually done with the *find()* method
- Consider an inventory collection created with the following *insertMany()* statements
 - Properties - item, quantity, size & status

```
> db.inventory.insertMany([
...   { item: "journal", qty: 25, size: { h: 14, w: 21, uom: "cm" }, status: "A" },
...   { item: "notebook", qty: 50, size: { h: 8.5, w: 11, uom: "in" }, status: "A" },
...   { item: "paper", qty: 100, size: { h: 8.5, w: 11, uom: "in" }, status: "D" },
...   { item: "planner", qty: 75, size: { h: 22.85, w: 30, uom: "cm" }, status: "D" },
...   { item: "postcard", qty: 45, size: { h: 10, w: 15.25, uom: "cm" }, status: "A" },
... ]);
{
  "acknowledged" : true,
  "insertedIds" : [
    ObjectId("5f98480f663182f750282297"),
    ObjectId("5f98480f663182f750282298"),
    ObjectId("5f98480f663182f750282299"),
    ObjectId("5f98480f663182f75028229a"),
    ObjectId("5f98480f663182f75028229b")
  ]
}
```

CRUD in Mongo

- A basic query using the find method takes a JSON-like query specification
- This example finds all documents with the status “D”
- Equivalent to the SQL at right in a relational database

```
> db.inventory.find( { status: "D" } )
{ "_id" : ObjectId("5f98480f663182f750282299"), "
" }
{ "_id" : ObjectId("5f98480f663182f75028229a"), "
"D" }
> [ ]
```

```
SELECT * FROM inventory WHERE status = "D"
```

CRUD in Mongo

- IN-style queries use the *\$in* keyword
- This query returns all documents whose status is in the given list
- Equivalent to the given SQL

```
db.inventory.find( { status: { $in: [ "A", "D" ] } } )
{ "_id" : ObjectId("5f98480f663182f750282297"), "item" : "j"
}
{ "_id" : ObjectId("5f98480f663182f750282298"), "item" : "n"
  "A" }
{ "_id" : ObjectId("5f98480f663182f750282299"), "item" : "p"
  "D" }
{ "_id" : ObjectId("5f98480f663182f75028229a"), "item" : "p"
  "D" }
{ "_id" : ObjectId("5f98480f663182f75028229b"), "item" : "p"
  "A" }
> []
```

```
SELECT * FROM inventory WHERE status in ("A", "D")
```

CRUD in Mongo

- AND queries are comma separated
- This query returns all documents with status "A" AND a quantity less than 30
 - Note the less than specification, using the *\$lt* keyword
 - Many comparison keywords list this are available: *\$eq*, *\$gt*, *\$gte*, *\$in*, *\$nin*, etc.

```
> db.inventory.find( { status: "A", qty: { $lt: 30 } } )
{ "_id" : ObjectId("5f98480f663182f750282297"), "item" :
"
}>
```

```
SELECT * FROM inventory WHERE status = "A" AND qty < 30
```

CRUD in Mongo

- OR queries use the *\$or* keyword, which takes an array of conditions
- Here we are finding all inventory documents with status “A” or quantity less than 30

```
carson@grim
> db.inventory.find( { $or: [ { status: "A" }, { qty: { $lt: 30 } } ] } )
{ "_id" : ObjectId("5f98480f663182f750282297"), "item" : "journal", "qty" : 25 }
{ "_id" : ObjectId("5f98480f663182f750282298"), "item" : "notebook", "qty" : 20 }
{ "_id" : ObjectId("5f98480f663182f75028229b"), "item" : "postcard", "qty" : 15 }
```

```
SELECT * FROM inventory WHERE status = "A" OR qty < 30
```

CRUD in Mongo

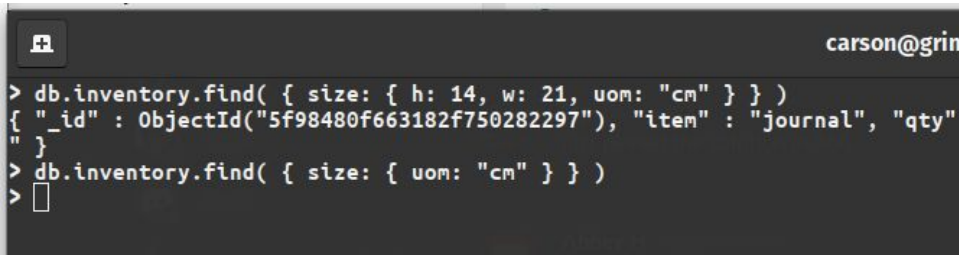
- ANDs and ORs can be combined with a nested query tree
- Use a standard comma separated JSON object for AND and the `$or:[]` syntax for ORs

```
carson
> db.inventory.find( {
...   status: "A",
...   $or: [ { qty: { $lt: 30 } }, { item: /^p/ } ]
... } )
{ "_id" : ObjectId("5f98480f663182f750282297"), "item" : "journal", "
" }
{ "_id" : ObjectId("5f98480f663182f75028229b"), "item" : "postcard",
: "A" }
> []
```

```
SELECT * FROM inventory WHERE status = "A" AND ( qty < 30 OR item LIKE "p%")
```

CRUD in Mongo

- What about querying embedded/nested data?
- Multiple syntaxes with slight differences
 - Obvious nested syntax requires an exact match on all fields
 - Cannot omit any fields, or a document won't match
 - Note that second query does not match any documents!

A terminal window with a dark background. The title bar shows a window icon and the text 'carson@grin'. The terminal contains three lines of MongoDB shell commands. The first command is a query for a document with a specific size and unit of measure, which returns a single document. The second command is a query for documents with a specific unit of measure, which returns an empty array. The third command is a prompt character '>'.

```
carson@grin
> db.inventory.find( { size: { h: 14, w: 21, uom: "cm" } } )
{ "_id" : ObjectId("5f98480f663182f750282297"), "item" : "journal", "qty" : 1 }
> db.inventory.find( { size: { uom: "cm" } } )
> []
```

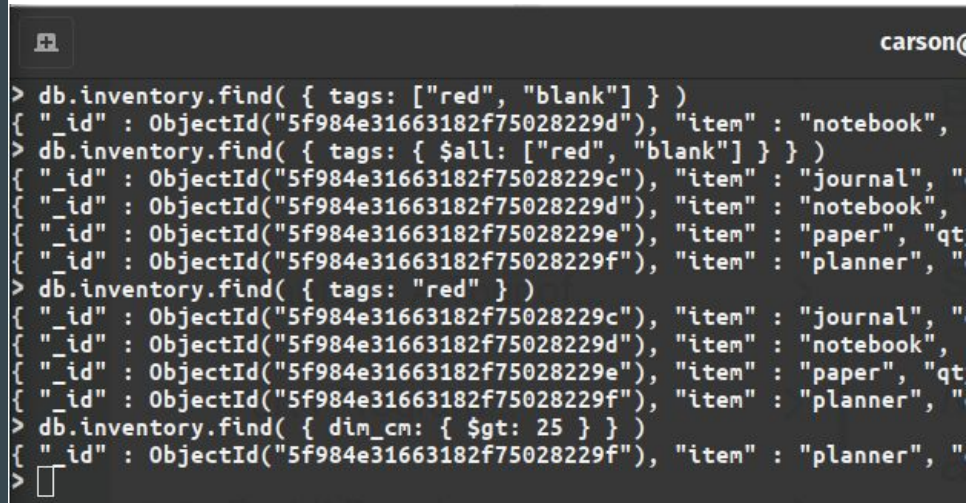

CRUD in Mongo

- More intuitive syntax is available if you use the dot syntax
- This syntax does *not* require a total match

```
carson
> db.inventory.find( { "size.uom": "cm" } )
{ "_id" : ObjectId("5f98480f663182f750282297"), "item" : "journal",
  "size" : "A" }
{ "_id" : ObjectId("5f98480f663182f75028229a"), "item" : "planner",
  "size" : "D" }
{ "_id" : ObjectId("5f98480f663182f75028229b"), "item" : "postcard",
  "size" : "A" }
>
```

CRUD in Mongo

- Arrays can be queried in various ways
 - Exact match
 - Unordered match
 - Contains
 - Conditions
- Nested documents can be queried against via the same dot syntax used for fields



A screenshot of a MongoDB shell terminal window. The window has a title bar with a plus icon and the name 'carson'. The terminal shows a series of commands and their corresponding JSON results. The commands use the `db.inventory.find()` method with various query filters. The results are JSON arrays of objects, each containing an `_id` (ObjectId) and an `item` string. The items include 'notebook', 'journal', 'paper', and 'planner'.

```
> db.inventory.find( { tags: ["red", "blank"] } )
{ "_id" : ObjectId("5f984e31663182f75028229d"), "item" : "notebook",
  ...
}
> db.inventory.find( { tags: { $all: ["red", "blank"] } } )
{ "_id" : ObjectId("5f984e31663182f75028229c"), "item" : "journal", "
  ...
}
{ "_id" : ObjectId("5f984e31663182f75028229d"), "item" : "notebook",
  ...
}
{ "_id" : ObjectId("5f984e31663182f75028229e"), "item" : "paper", "qt
  ...
}
{ "_id" : ObjectId("5f984e31663182f75028229f"), "item" : "planner", "
  ...
}
> db.inventory.find( { tags: "red" } )
{ "_id" : ObjectId("5f984e31663182f75028229c"), "item" : "journal", "
  ...
}
{ "_id" : ObjectId("5f984e31663182f75028229d"), "item" : "notebook",
  ...
}
{ "_id" : ObjectId("5f984e31663182f75028229e"), "item" : "paper", "qt
  ...
}
{ "_id" : ObjectId("5f984e31663182f75028229f"), "item" : "planner", "
  ...
}
> db.inventory.find( { dim_cm: { $gt: 25 } } )
{ "_id" : ObjectId("5f984e31663182f75028229f"), "item" : "planner", "
  ...
}
>
```

CRUD in Mongo

- Field Selection: to select specific fields from a document, you pass in a second JSON object
 - Passing a 1 for a field name indicates it is included
 - Passing 0 for a field name indicates all other *except* this field should be returned
 - Can use dot syntax to include or exclude embedded documents

```
carson@grimlock: ~  
> db.inventory.find( { status: "A" }, { item: 1, status: 1 } )  
{ "_id" : ObjectId("5f98480f663182f750282297"), "item" : "journal", "status" : "A" }  
{ "_id" : ObjectId("5f98480f663182f750282298"), "item" : "notebook", "status" : "A" }  
{ "_id" : ObjectId("5f98480f663182f75028229b"), "item" : "postcard", "status" : "A" }  
>
```

```
SELECT _id, item, status from inventory WHERE status = "A"
```

CRUD in Mongo

- Updates are done with one of the following methods:

- `db.collection.updateOne(<filter>, <update>, <options>)`
- `db.collection.updateMany(<filter>, <update>, <options>)`
- `db.collection.replaceOne(<filter>, <update>, <options>)`
- `db.collection.update(<filter>, <update>, <options>)`

```
db.inventory.updateOne(
  { item: "paper" },
  {
    $set: { "size.uom": "cm", status: "P" },
    $currentDate: { lastModified: true }
  }
)
```

CRUD in Mongo

- `updateOne()` takes a condition and an update expression and updates the *first match*
 - Here we update the first inventory document whose item is “paper”
 - sets the embedded document *size uom* property to “cm”
 - Sets the status to “P”
 - Sets lastModified to the current date

```
db.inventory.updateOne(
  { item: "paper" },
  {
    $set: { "size.uom": "cm", status: "P" },
    $currentDate: { lastModified: true }
  }
)
```

CRUD in Mongo

- *updateMany()* takes a condition and an update expression and updates all matches
 - Here we update all inventory documents whose item is “paper”
 - sets the embedded document *size uom* property to “in”
 - Sets the status to “P”
 - Sets lastModified to the current date

```
db.inventory.updateMany(  
  { "qty": { $lt: 50 } },  
  {  
    $set: { "size.uom": "in", status: "P" },  
    $currentDate: { lastModified: true }  
  }  
)
```

CRUD in Mongo

- *replaceOne()* takes a condition and an update expression and *replaces* all matches
 - This will replace the entire document, not just update fields

```
db.inventory.replaceOne(  
  { item: "paper" },  
  { item: "paper", instock: [ { warehouse: "A", qty: 60 } ],  
  )
```

CRUD in Mongo

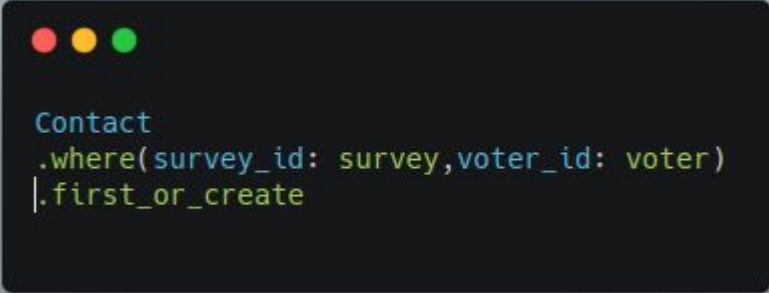
- Upsert Support

- Any update operation can be converted into an *upsert* by including the `{upsert : true}` option
- If no match, the document is inserted instead
- Here, if no product with id 6 exists, the data will be inserted instead

```
db.products.updateMany(  
  { _id: 6 },  
  { $set: {price: 999} },  
  { upsert: true}  
)
```

CRUD in Mongo

- Upserts are very useful in many online systems
 - Rails ORM ActiveRecord supports a similar pattern with *first_or_create()*

A terminal window with a dark background and three colored window control buttons (red, yellow, green) in the top-left corner. It contains Ruby code for ActiveRecord.

```
Contact  
.where(survey_id: survey, voter_id: voter)  
|.first_or_create
```

CRUD in Mongo

- Delete operations

- `db.collection.deleteOne(<filter>, <options>)`
- `db.collection.deleteMany(<filter>, <options>)`
 - *<Tell Funny OR story here>*

```
db.inventory.deleteMany({ status : "A" })
```

Text Search

- Mongo supports broad, google-like text search out of the box
 - First example: general text search
 - Second example: exact match for “coffee shop”
 - Third excludes “coffee”
- Excellent functionality compared with most RDBMS
 - Especially a decade ago!

```
db.stores.find( { $text: { $search: "java coffee shop" } } )
```

```
db.stores.find( { $text: { $search: "\"coffee shop\"" } } )
```

```
db.stores.find( { $text: { $search: "java shop -coffee" } } )
```

Indexing

- Mongo supports the easy creation of indexes
 - Supports geospatial indexing, text indexing, multi-key indexing, etc.
- Index usage can be viewed via the *\$indexStats* keyword

```
db.collection.createIndex( { name: -1 } )
```

Transactions

- Historically Mongo has not had a good reputation for data persistence
 - Arguments over whether this is FUD or not
 - *I would use an RDBMs for crucial data, myself*
- Mongo now has a transaction API
 - API is based on distributed transactions, pretty complex

```
# Step 2: Start a client session.
with client.start_session() as session:
    # Step 3: Use with_transaction to start a transaction
    session.with_transaction(
        callback, read_concern=ReadConcern('local'),
        write_concern=wc_majority,
        read_preference=ReadPreference.PRIMARY)
```

When To Use Mongo?

- As much of a technology grump as I am, I think Mongo has a place in many systems
- Useful for:
 - Non-core data
 - Data streams
 - Flexible, early data modeling
 - Data that “doesn’t matter”
- I would still recommend RDBMS for core, “must be correct” data

MongoDB	RDBMS
Database	Database
Collection	Table
Document	Row
Field	Column

MongoDB

- MongoDB is a Document Database
 - In contrast with a *Relational Database*
 - Optimized for the storage and retrieval of documents
- We reviewed many of the core operations for Mongo
 - CRUD
 - Indexes
- A good SQL/Mongo mapping document
 - <https://docs.mongodb.com/manual/reference/sql-comparison/>
- Next time we will talk about the aggregation pipeline in Mongo



MONTANA
STATE UNIVERSITY