



MONTANA
STATE UNIVERSITY

Database Systems

...

History & Context

What is a Database?

- A database is an organized collection of structured information, or data, typically stored electronically in a computer system
- A database is usually controlled by a database management system (DBMS)

TrackId	Name	AlbumId	MediaTypeId	GenreId	Composer	MusicianId
1	1 For Those About To...	1	1	1	1 Angus Young, Malco...	
2	2 Balls to the Wall	2	2	2	1 <null>	
3	3 Fast As a Shark	3	2	2	1 F. Baltes, S. Kauf...	
4	4 Restless and Wild	3	2	2	1 F. Baltes, R.A. Sm...	
5	5 Princess of the Da...	3	2	2	1 Deaffy & R.A. Smit...	
6	6 Put The Finger On ...	1	1	1	1 Angus Young, Malco...	
7	7 Let's Get It Up	1	1	1	1 Angus Young, Malco...	
8	8 Inject The Venom	1	1	1	1 Angus Young, Malco...	
9	9 Snowballed	1	1	1	1 Angus Young, Malco...	
10	10 Evil Walks	1	1	1	1 Angus Young, Malco...	
11	11 C.O.D.	1	1	1	1 Angus Young, Malco...	
12	12 Breaking The Rules	1	1	1	1 Angus Young, Malco...	
13	13 Night Of The Long ...	1	1	1	1 Angus Young, Malco...	
14	14 Spellbound	1	1	1	1 Angus Young, Malco...	
15	15 Go Down	4	1	1	1 AC/DC	
16	16 Dog Eat Dog	4	1	1	1 AC/DC	
17	17 Let There Be Rock	4	1	1	1 AC/DC	
18	18 Bad Boy Boogie	4	1	1	1 AC/DC	
19	19 Problem Child	4	1	1	1 AC/DC	
20	20 Overdose	4	1	1	1 AC/DC	
21	21 Hell Ain't A Bad P...	4	1	1	1 AC/DC	
22	22 Whole Lotta Rosie	4	1	1	1 AC/DC	
23	23 Walk On Water	5	1	1	1 Steven Tyler, Joe ...	
24	24 Love In An Elevator	5	1	1	1 Steven Tyler, Joe ...	


DBMS

- Provides security
- Handles schema modification
- Provides isolation
- Supports a multi-user environment that allows parallel data access and manipulation

TrackId	Name	AlbumId	MediaTypeId	GenreId	Composer	Milli
1	1 For Those About To...	1	1	1	1 Angus Young, Malco...	
2	2 Balls to the Wall	2	2	1	1 <null>	
3	3 Fast As a Shark	3	2	1	1 F. Baltes, S. Kauf...	
4	4 Restless and Wild	3	2	1	1 F. Baltes, R.A. Sm...	
5	5 Princess of the Da...	3	2	1	1 Deaffy & R.A. Smit...	
6	6 Put The Finger On ...	1	1	1	1 Angus Young, Malco...	
7	7 Let's Get It Up	1	1	1	1 Angus Young, Malco...	
8	8 Inject The Venom	1	1	1	1 Angus Young, Malco...	
9	9 Snowballed	1	1	1	1 Angus Young, Malco...	
10	10 Evil Walks	1	1	1	1 Angus Young, Malco...	
11	11 C.O.D.	1	1	1	1 Angus Young, Malco...	
12	12 Breaking The Rules	1	1	1	1 Angus Young, Malco...	
13	13 Night Of The Long ...	1	1	1	1 Angus Young, Malco...	
14	14 Spellbound	1	1	1	1 Angus Young, Malco...	
15	15 Go Down	4	1	1	1 AC/DC	
16	16 Dog Eat Dog	4	1	1	1 AC/DC	
17	17 Let There Be Rock	4	1	1	1 AC/DC	
18	18 Bad Boy Boogie	4	1	1	1 AC/DC	
19	19 Problem Child	4	1	1	1 AC/DC	
20	20 Overdose	4	1	1	1 AC/DC	
21	21 Hell Ain't A Bad P...	4	1	1	1 AC/DC	
22	22 Whole Lotta Rosie	4	1	1	1 AC/DC	
23	23 Walk On Water	5	1	1	1 Steven Tyler, Joe ...	
24	24 Love In An Elevator	5	1	1	1 Steven Tyler, Joe ...	

History

- 1960s - The first databases began to appear
- Databases evolved from flat files
- IBM SABRE system for airlines
 - Helped American Airlines handle reservations



History


- 1970s - E.F. Cobb proposes the relational database model
- Ingres
 - QUEL query language
- System R
 - SEQUEL query language

INGRES




History

- 1970s - E/R Diagrams are proposed
- Application designers could abstract away from table design



History

- 1970s - E/R Diagrams are proposed
- Application designers could abstract away from table design...
 - *Is this a good thing?*



History

- 1980s to early 1990s:
The Golden Age of Databases
- SQL becomes a standard
- Many new players
- *Lots* of money was made



History

- Early 1990s: Big Shakeout
- Many vendors went bankrupt or became bit players
- Even **Oracle** came close to bankruptcy in 1990
 - Accounting shenanigans
 - "an incredible business mistake"




History

- Late 1990s: Internet Boom
- A resurgence in vendors, Oracle and DB2 dominant
- Microsoft started making inroads with a fork of Sybase called SQLServer
- Today these are the three dominant commercial players



History

- Late 1990s: Open Source Databases
- 1995 MySQL
 - Bought by oracle
- 1996 Postgresql
 - Based on Ingres at Berkeley
 - Go Bears
- 2000 SQLite
 - Super lightweight database
 - We are going to use it!



History

- Post 2000 - Maturity
- Three major commercial vendors
- Two major open source options
- Rise of NoSQL
 - We will look at two
 - MongoDB
 - Redis














Today

- Today's Database Rankings:

<https://db-engines.com/en/ranking>

- Rankings are pretty stable
- Oracle continues to lead
- SQLite is becoming more popular

Rank			DBMS
Jul 2020	Jun 2020	Jul 2019	
1.	1.	1.	Oracle 
2.	2.	2.	MySQL 
3.	3.	3.	Microsoft SQL Server 
4.	4.	4.	PostgreSQL 
5.	5.	5.	MongoDB 
6.	6.	6.	IBM Db2 
7.	7.	7.	Elasticsearch 
8.	8.	8.	Redis 
9.	9.	 11.	SQLite 
10.	10.	10.	Cassandra 

What is a Database?

- A system that stores and retrieves data
- A query processor
- Some sort of data definition language (DDL)
- Typically provides network access
 - Not always! SQLite doesn't!
- Makes some guarantees about data




Sidebar: Why A Cylinder?

- Probably due to the original appearance of hard disks



Sidebar: Why A Cylinder?

- Probably due to the original appearance of hard disks



Sidebar: Why A Cylinder?

- Probably due to the original appearance of hard disks



Sidebar: Why A Cylinder?

- Probably due to the original appearance of hard disks
- Oracle Headquarters...



Anyway, So, Back to Data

- DBMS Guarantees = Transactions (mostly)
- ACID
 - Atomicity - All completes or all fails
 - Consistency - Data consistency constraints are enforced
 - Isolation - Transactions complete as if no other transaction occurred
 - Durability - Data won't be lost



ACID, Who needs it?

- NoSQL stores often relax or more of these constraints for performance reasons
- Very popular in the early 2010s!
- Since then the industry has shown some regrets
- A return to standard SQL systems or mixed systems



End of Course Goals

- Be comfortable with the following:
 - Designing database schemas
 - Writing queries
 - Doing so from the web
 - Tuning queries for performance
 - Discussing general database architecture in a local and cloud context
- Be employable





Vlad Mihalcea

@vlad_mihalcea



Some companies have tech interviews that validate whether you can send a rocket to the moon, when, in fact, their day-to-day applications just fetch data from a database and display it on the UI.

9:20 PM · Jul 20, 2020 · [Twitter for Android](#)

232 Retweets and comments **1K** Likes



MONTANA
STATE UNIVERSITY