



MONTANA
STATE UNIVERSITY


Database Design

...

Designing a Proper Database

Designing A Database


- We have worked with E/R diagrams
- Often a good *first pass* at a database design
- Makes your needs concrete
 - May expose data problems
 - May give you insights into what data is needed
 - N-to-n → Entity transformation is often valuable



Designing A Database

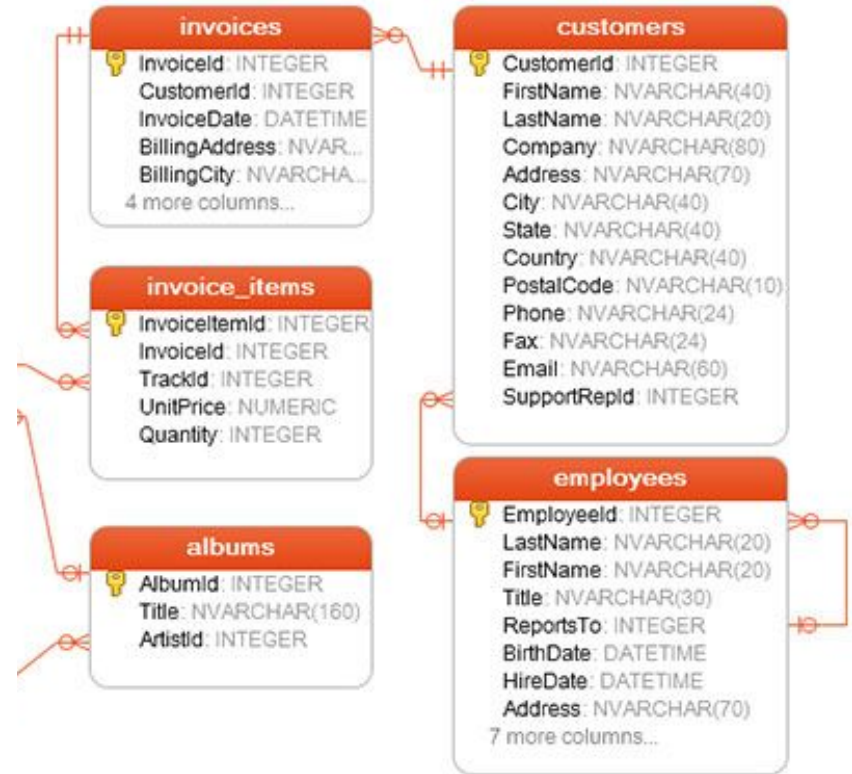
- General Process

- What are the core nouns in my system?
- What are their relationships?
 - 1-1
 - 1-n
 - N-n (hmm)
- What other nouns are needed to properly model things?
 - Things I haven't thought of
 - Decomposing existing nouns
- GOTO step 1



Designing A Database

- Once you have an E/R diagram you can transform it into a relational schema
- A relational schema typically reveals more detailed needs in your data model




Designing A Database

- Finally, this can be transformed into actual database definition statements
- Yet more details and issues will come up
- (we cover DDL later in the course)

```
CREATE TABLE Person(  
    PersonID INT IDENTITY (1,1) CONSTRAINT PK_PersonID PRIMARY KEY,  
    FirstName NVARCHAR(20),  
    LastName NVARCHAR(25)  
);
```

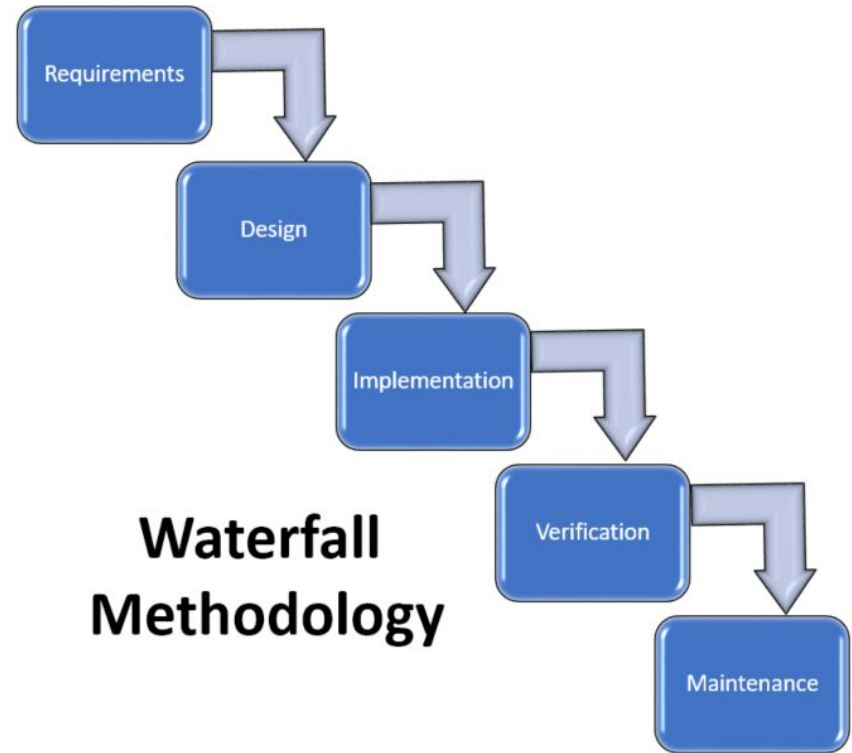
E/R to Database Translation

- Entity → Table
- Relation → FK(s)
 - 1-1 - pick the “lesser” relation to store the fk on
 - 1-n - store the fk on the N side of the relation
 - N-n - Is there an entity here?
 - If yes, add it
 - If no, create a join table (a table with two FKs only)



Designing A Database

- How realistic is this?
 - E/R →
Relational Schema →
DDL?
- Typically... not very realistic



Designing A Database

"No plan survives contact with the enemy."

--Helmuth von Moltke the Elder




Designing A Database

- As you begin working with the database schema, problems will inevitable come up
- You will address them in the database directly
- The E/R diagrams will become increasingly out of date

Application Architecture

- **End-user view**
 - Performance
 - Availability
 - Usability
 - Security
- **Business view**
 - Targeted market
 - Time to market
 - Cost and benefits
 - Projected time line
 - Integration with legacy systems
- **Developer view**
 - Maintainability
 - Portability
 - Reusability
 - Testability



See ISO/IEC 9126-2001

The image is published by Microsoft Corp.

Enterprise Software Application Architecture




Designing A Database

- Application architects can become increasingly irrelevant if they continue to use E/R Diagrams and other artifacts to discuss the system
 - They may also become increasingly annoying
- Maintain a sense of humor during this process

Application Architecture

- **End-user view**
 - Performance
 - Availability
 - Usability
 - Security
- **Business view**
 - Targeted market
 - Time to market
 - Cost and benefits
 - Projected time line
 - Integration with legacy systems
- **Developer view**
 - Maintainability
 - Portability
 - Reusability
 - Testability



See ISO/IEC 9126-2001

The image is published by Microsoft Corp.

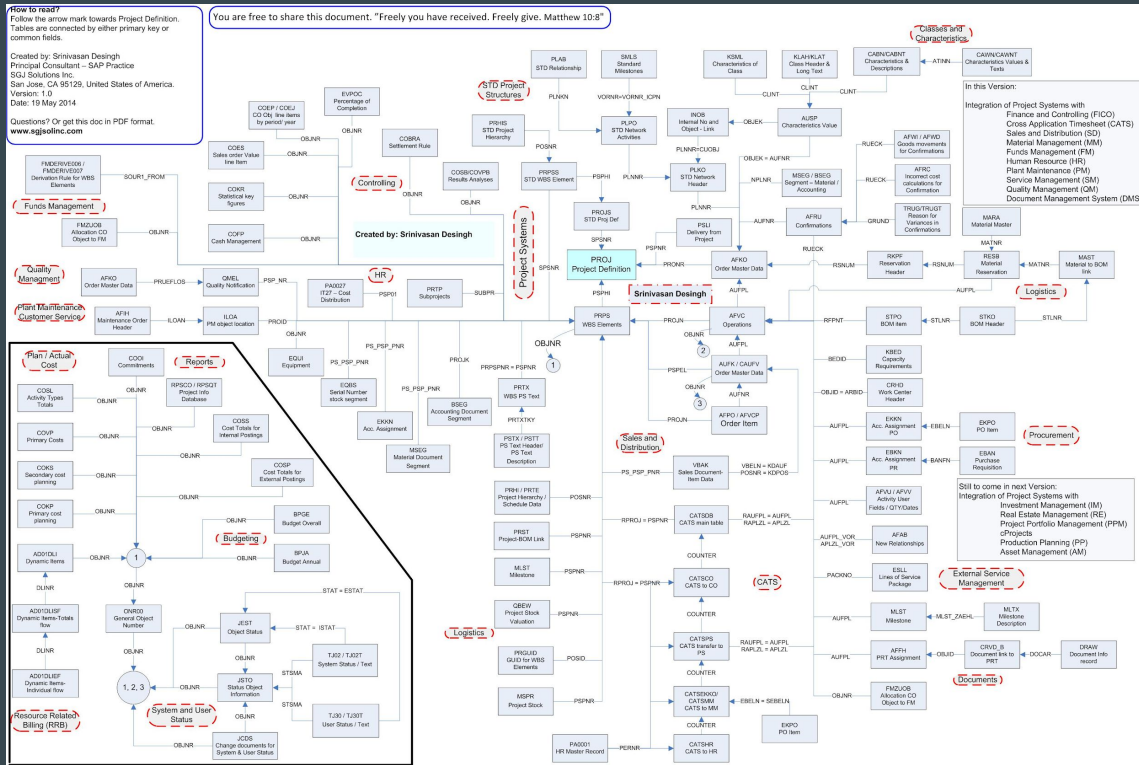
Enterprise Software Application Architecture



Case Study - SAP

- SAP is one of the most valuable software companies in the world
- Worth 2 billion dollars today
- Sells software to help big companies manage things like inventory
- Here is a small part of their schema...

Case Study - SAP



Case Study - SAP

- The SAP schema is famously *a nightmare*
 - It is **huge**
 - **More than 8000 tables**
 - Large parts of it make no sense
 - The table names are often in german
 - There are normalization issues all over it
 - We will talk about normalization in the next lecture
 - **SAP is still worth 2 billion dollars**

Advice On Database Modeling

- As the system matures, accept that data modeling mistakes made early on are going to stick with you
- The database schema get harder to fix over time
 - That's OK! Learn to live with it.
- QUESTION: Why is a database so much harder to fix than code?
- LeadDyno DB Story Time!

Advice On Database Modeling

- Start with E/R diagramming
 - But don't take it too seriously
- Get to a real world database **ASAP**
- Iterate your app and database as quickly as possible early on when database size is relatively small
- Try to figure out the crucial entities in your system and get those as right as possible
- Read up on stoicism

Polymorphism & Databases

...


A Mismatch?

Coding With Databases

- Almost all databases are fronted by some software
- That software is written in a particular language
- That language is *probably* object oriented
 - Java (this class)
 - Javascript :(
 - PHP XD


Polymorphism

- You are probably familiar with this idea from your Object Oriented classes
 - Super-classes
 - Sub-classes
 - Sub-classes extend the super-class
 - Add methods and *attributes*




The Relational Model

- Relations are just tables and foreign keys
- How should we model this object hierarchy if we want to store it in a table?



Three Approaches

- Single Table Inheritance
- Class Table Inheritance
- Concrete Table Inheritance



Single Table Inheritance

- A single table is used for all sub-class instances
- The columns are the union of all columns of sub-classes
- Advantages?
- Disadvantages?

Single Table Inheritance

Shape Type	x	y	radius	length	width	height
Triangle	1	5			10	20
Rectangle	3	6		20	10	
Circle	8	3	6			
Triangle	4	4			6	9

Class Table Inheritance

- There is one table per class in the object hierarchy
- Sub-classes include a foreign key reference to their parent classes
- Advantages?
- Disadvantages?

Shapes

id	x	y
1	1	5
2	3	6
3	8	3
4	4	4

Circles

shape_id	radius
3	6

Rectangles

shape_id	length	width
2	20	10

Triangles

shape_id	length	width
1	10	20
4	6	9

Concrete Table Inheritance

- There is one table per **concrete** class in your object hierarchy
- Advantages?
- Disadvantages?

Circles

x	y	radius
8	3	6

Rectangles

x	y	length	width
3	6	20	10

Triangles

x	y	width	height
1	5	10	20
4	4	6	9

Polymorphism In Practice

- Most systems that I have experience with use a mix:
 - Single Table Inheritance for closely related things with many foreign keys in common
 - Concrete Table Inheritance for more distant relations
 - I have never seen class table inheritance work out well
- My advice is to focus on foreign keys: *the more keys two objects have in common, the more likely you are to prefer single table inheritance*

What Are We Dealing With Here?

- This is a single instance of a more general problem...

The Object-Relational Impedance Mismatch

“The object-relational impedance mismatch is a set of conceptual and technical difficulties that are often encountered when a relational database management system (RDBMS) is being served by an application program (or multiple application programs) written in an object-oriented programming language or style, particularly because objects or class definitions must be mapped to database tables defined by a relational schema.”

The Object-Relational Impedance Mismatch

- While inheritance is one issue, there are many others:
 - Encapsulation isn't part of the relational model
 - Interfaces don't exist at the relational level
 - Field accessibility isn't specified at the relational level
 - Database transactions do not map well to objects
 - And so on
- We will discuss this more thoroughly when we talk about Object/Relational mapping tools

The Object-Relational Impedance Mismatch

- There is, however, a fundamental conceptual and cultural difference between Object Oriented and Relational thinking
- Object Oriented
 - Imperative
 - Nouns + Verbs
 - Graphs
- Relational Thinking
 - Declarative
 - Normalized
 - Set relationships

Object Databases

- Some people reject the relational model entirely
- OODBMS (Object-oriented database management systems) explicitly encode OO semantics
- Smalltalk/GemTalk
- Java/db4o
- A small but *passionate* community





MONTANA
STATE UNIVERSITY